



Figure 1. Sensitivity of ranking performance to κ . Test AUTOC (mean $\pm$ s.e. over five seeds) as a function of the smoothing parameter κ for different training sizes. *Higher is better.* Our *Rank-Learner* is relatively insensitive to the choice of κ over the considered range. At small n , *Rank-Learner* consistently outperforms the plug-in ranker across the full κ range. As n increases, performance differences become smaller, consistent with improved nuisance estimation and convergence toward oracle performance.