

Table 2. **Synthetic benchmark (additional baselines)**. Test AUTOC and average policy value (mean $\pm$ std dev over five seeds) across training sizes. *Higher is better*. The oracle column reports the metrics obtained by ranking the test set using the true treatment effects. Best mean based on unrounded values is shown in **bold**, second best underlined.

Metric	Method	$n = 500$	$n = 1,000$	$n = 2,000$	oracle
AUTOC ($\uparrow$)	Tree ranker	1.27 ± 0.03	1.31 ± 0.01	1.34 ± 0.01	1.40
	DR plug-in ranker	1.28 ± 0.01	1.32 ± 0.02	1.35 ± 0.01	
	PairNet + sort	1.29 ± 0.01	<u>1.33 ± 0.01</u>	<u>1.35 ± 0.01</u>	
	FlexTENet + sort	1.30 ± 0.02	1.32 ± 0.01	1.34 ± 0.02	
	Rank-learner (<i>ours</i>)	<u>1.30 ± 0.03</u>	1.34 ± 0.01	1.37 ± 0.00	
$\bar{V}$ ($\uparrow$)	Tree ranker	0.55 ± 0.01	0.57 ± 0.00	0.57 ± 0.00	0.59
	DR plug-in ranker	0.55 ± 0.00	0.56 ± 0.01	0.58 ± 0.00	
	PairNet + sort	0.55 ± 0.01	<u>0.57 ± 0.00</u>	<u>0.58 ± 0.00</u>	
	FlexTENet + sort	<u>0.56 ± 0.00</u>	0.57 ± 0.00	0.57 ± 0.01	
	Rank-learner (<i>ours</i>)	0.56 ± 0.01	0.57 ± 0.00	0.58 ± 0.00	