

Table 2. **Synthetic benchmark (additional baselines).** Test AUTO C and average policy value (mean $\pm$ std dev over five seeds) across training sizes. *Higher is better.* The oracle column reports the metrics obtained by ranking the test set using the true treatment effects. Best mean based on unrounded values is shown in **bold**.

Metric	Method	$n = 250$	$n = 500$	$n = 1,000$	$n = 2,000$	oracle
AUTO C ($\uparrow$)	Tree ranker	1.20 ± 0.04	1.27 ± 0.03	1.31 ± 0.01	1.34 ± 0.01	1.40
	DR plug-in ranker	1.20 ± 0.09	1.28 ± 0.01	1.32 ± 0.02	1.35 ± 0.01	
	FlexTENet + sort	1.23 ± 0.06	1.30 ± 0.02	1.33 ± 0.01	1.34 ± 0.02	
	Rank-learner (<i>ours</i>)	1.28 ± 0.03	1.30 ± 0.03	1.34 ± 0.01	1.37 ± 0.00	
$\bar{V}$ ($\uparrow$)	Tree ranker	0.53 ± 0.01	0.55 ± 0.01	0.57 ± 0.00	0.57 ± 0.00	0.59
	DR plug-in ranker	0.53 ± 0.03	0.55 ± 0.00	0.56 ± 0.01	0.58 ± 0.00	
	FlexTENet + sort	0.54 ± 0.01	0.56 ± 0.00	0.57 ± 0.00	0.57 ± 0.01	
	Rank-learner (<i>ours</i>)	0.55 ± 0.01	0.56 ± 0.01	0.57 ± 0.00	0.58 ± 0.00	