
Table 3. Synthetic benchmark (nuisance misspecification): Test AUTO C (mean $\pm$ std dev over five seeds) for $n = 250$ under the standard setting and under deliberately misspecified linear nuisance models. *Higher is better.* The oracle column reports the AUTO C obtained by ranking the test set using the true treatment effects. Best mean in each column is shown in **bold**.

Method	Standard nuisances	Linear nuisances	oracle
T-learner	0.96 ± 0.14	0.76 ± 0.07	1.40
DR-learner	1.16 ± 0.12	1.09 ± 0.11	
Plug-in ranker	0.95 ± 0.14	0.66 ± 0.06	
Rank-learner (ours)	1.28 ± 0.03	1.19 ± 0.04	
